

## Deep Learning Technical Report: Hybrid Multimodal Classification

**Project Goal:** To build and evaluate a hybrid deep learning model capable of classifying patient status (target variable: sport\_ability) using a combination of **12-lead Electrocardiogram (ECG) time-series data** and **structured clinical/anthropometric tabular data**.

### 1. Pipeline Overview and Data Handling

The pipeline is composed of four main stages: Data Loading, Multimodal Preprocessing, Model Training, and K-Fold Cross-Validation. A core feature of this solution is its ability to process two heterogeneous data types simultaneously.

#### 1.1 Data Loading and ECG Processing

The script is specifically configured to load **real data** from .mat (MATLAB) files, which contain the raw 12-lead ECG signals, and tabular data from CSV files.

#### A. Signal Preprocessing


Raw ECG signals are often noisy and contain baseline wander. The following digital signal processing steps are applied to each of the 12 leads for every subject:

1. **Baseline Correction:** The mean of the signal is subtracted to center it around zero, removing DC offsets.
2. **Bandpass Filtering:** A Butterworth Bandpass Filter (1.0 Hz to 40.0 Hz) is applied to remove low-frequency baseline drift and high-frequency muscle noise.
3. **Notch Filtering:** A 50Hz Notch Filter is used to eliminate powerline interference, a common artifact in clinical ECG recordings.

#### B. Feature Extraction (Beat Segmentation)

Instead of feeding the entire, long ECG time series (5000 samples) into the model, a representative heartbeat is extracted:

1. **R-Peak Detection:** The **R-peak**, the most prominent point of the QRS complex, is detected, typically on Lead II, using a signal differentiation and integration technique to enhance peak visibility.



2. **Beat Segmentation:** Fixed-length segments (600 samples, centered around the R-peak) are extracted for every detected beat.

3. **Averaging:** All segmented beats for a subject are **averaged** to create a single, high-fidelity, representative heartbeat morphology (12 leads x 600 samples). This drastically reduces dimensionality while preserving the key diagnostic morphology (P-QRS-T complex).
4. **Z-Score Normalization:** The final averaged beat is normalized per-lead.

## 1.2 Tabular Data Preprocessing

The tabular data—composed of clinical and demographic variables—undergoes a structured preprocessing workflow:

1. **Feature Engineering:** New, informative features are calculated from the existing ones:
  - o **Body Mass Index (BMI)**
  - o **Body Surface Area (BSA)**
2. **Imputation:** An **IterativeImputer** (using a technique like Bayesian Ridge Regression) is fit *only* on the training data to estimate and fill missing values (NaNs) in the numerical features.
3. **Scaling and Encoding:**
  - o **Numerical Features** (age, weight, BMI, etc.) are scaled using a **StandardScaler** to have zero mean and unit variance.
  - o **Categorical Features** (sex, sport\_classification) are converted to numerical format using **OneHotEncoder**.
4. **Output:** The final tabular input is a vector of 9 features.

## 2. Model Architecture: The Three-Branch CNN

The network is a **Three-Branch Convolutional Neural Network (CNN)** designed for multimodal fusion:

1. **ECG Frontal Branch:** A CNN processes the 6 frontal leads (I, II, III, aVR, aVL, aVF).
  2. **ECG Precordial Branch:** A CNN processes the 6 precordial leads (V1-V6).
  3. **Tabular Branch:** A small **Multi-Layer Perceptron (MLP)** processes the 9 tabular features.
- **Feature Extraction:** Both ECG branches use identical CNN blocks (Conv1D, ReLU, MaxPool) to independently learn spatial and temporal patterns from the two main physiological views of the heart.
  - **Feature Fusion:** The learned features from the three branches are **concatenated** (joined end-to-end) into a single, high-dimensional vector of 4,640 features.
  - **Classification:** A final fully connected head (with **Dropout** for regularization and **Sigmoid** activation for binary classification) takes the fused vector and outputs the final probability of the positive class.

### 3. Evaluation Methodology: K-Fold Cross-Validation

#### Why K-Fold?

The **Stratified K-Fold Cross-Validation** approach ( $K=5$ ) was chosen for rigorous evaluation:

- **Robustness:** It provides a much more reliable estimate of the model's true generalization performance than a single train/test split.
- **Full Data Utilization:** Every data point is used exactly once in the validation set and  $K-1$  times in the training set, which is crucial for smaller datasets.
- **Stratification:** Ensures that the proportion of positive and negative classes is preserved in every fold, preventing bias in the evaluation due to class imbalance in one specific split.

### 4. Analysis of Final Results

The model was trained for **10 epochs** across 5 folds. The final results reveal a model with decent discriminative power but a significant bias caused by **overfitting**.

#### 4.1 Aggregate Metrics Summary

Metric	Mean ( $\mu$ )	Standard Deviation ( $\sigma$ )	Interpretation
Mean AUC	0.7093	0.0821	The model has a fair ability to distinguish between the classes. The high $\sigma$ shows high variability between folds, indicating the model's sensitivity to the training data.
Mean Accuracy	0.7243	0.0639	The model correctly predicts the class for about 72.4% of subjects on average.
Mean Sensitivity	0.8694	0.0692	<b>Excellent</b> at detecting <b>True Positives</b> (i.e., when the subject <i>should</i> be classified as Positive, the model is correct 87% of the time).
Mean Specificity	0.4102	0.1016	<b>Poor</b> at detecting <b>True Negatives</b> (i.e., when the subject <i>should</i> be classified as Negative, the model is correct only 41% of the time).

#### 4.2 Overfitting and Bias Analysis

The primary issue revealed by the loss curves and the metrics is **overfitting** coupled with a strong **classification bias**:


1. **Overfitting Confirmation:** In Folds 3, 4, and 5, the **Training Loss** drops consistently (down to *approx* 0.3), while the **Validation Loss** stops decreasing around Epoch 4-5 and begins to rise significantly (up to *approx* 0.84). This

divergence is the hallmark of overfitting. The model has learned the training data noise but fails to generalize.

2. **Classification Bias:** The massive gap between **Sensitivity (approx 87%)** and **Specificity (approx 41%)** indicates the model has learned to predominantly predict the **Positive class**. It minimizes false negatives at the expense of creating many **False Positives**. This behavior is typical when the target variable is imbalanced, and the model minimizes the loss by favoring the more frequent class.

#### 4.3 Visual Results

The model results are visualized below, confirming the aggregate metrics and cross-fold variability.



#### **4.4 Recommendations for Next Steps**

To improve generalization and reduce the classification bias, the following steps are recommended:

1. **Early Stopping:** Implement a rule to stop training when the **Validation Loss** fails to decrease for several epochs. This is the most effective way to combat the observed overfitting.
2. **Addressing Imbalance:** Employ a **weighted Binary Cross-Entropy (BCE) loss function** to place a higher penalty on misclassifying the minority class, forcing the model to improve its Specificity.
3. **Regularization Tuning:** Increase the Dropout rate or introduce **L2 regularization (weight decay)** to the optimizer to further penalize complex models.